Ideal Customer Profile

A framework for prospecting and qualification

June 2026

THE THESIS

FireHydrant fits one buyer above all: **a software-driven business where downtime carries a direct, measurable cost** — held by an engineering organisation mature enough to require codified incident process, and already operating a modern observability stack. Where those conditions hold, the value case writes itself. Where they do not, no amount of selling will create durable demand.

The parameters that follow are ordered by how strongly each predicts fit. They resolve into two practical questions, addressed in turn: *whom to pursue* in prospecting, and *how to confirm* a genuine opportunity in qualification.

I. How to use this framework

Step one — Prospecting. Build pipeline against the signals observable from the outside: vertical and business model, company and engineering-team size, the presence of dedicated reliability roles, the observability stack, the collaboration platform, and any active buying trigger. These can be assessed before a single conversation.

Step two — Qualification. In conversation, confirm what only dialogue reveals. Treat the four Tier 1 criteria as a gate; then score the fit and winnability parameters on the matrix in Section III.

II. The parameters

TIER 1 Non-negotiable — disqualify if absent

Cost of downtime is measurable. Transactional or real-time revenue, contractual SLA penalties, or reliability sold as the product itself. Absent an articulable cost-per-minute or SLA exposure, there is no return-on-investment case.

Software-driven business. The core product or operation is an always-on digital service. Non-digital businesses, and pure IT-helpdesk buyers, fall outside the profile; the latter belongs to the Freshservice motion.

Dedicated reliability function. An SRE, Platform, or DevOps team owns production incidents. Where no such function exists, the pain is rarely acute enough to compel a purchase.

Modern observability stack. FireHydrant is a coordination layer atop monitoring, not a detection tool. Look for Datadog, Grafana, Prometheus, New Relic, or Honeycomb in active use.

TIER 2 Strong fit signals — the basis of the fit score

Process maturity and buying trigger. An organisation past the "move fast and break things" stage. The strongest catalysts are a recent painful or public outage, or a newly appointed Head of SRE or VP of Engineering mandated to formalise operations.

Engineering organisation shape. Specialised reliability roles, such that process cannot remain tribal: dedicated SREs and Platform engineers, formal on-call across multiple teams, and an incident-commander or reliability-lead role in place or being created. The number of product teams shipping to production is a better signal than raw headcount.

Company and engineering-team size. Outer bounds of 500 to 10,000 employees; within that range, the engineering team determines fit. Thirty to two hundred engineers represents the sweet spot.

Architecture complexity. A complex or microservice estate, where the service catalogue's ownership and dependency mapping earns its keep. Monolithic environments derive less from this differentiator.

Legacy alerting to displace. PagerDuty or Opsgenie already deployed. This is a trigger rather than a barrier; Opsgenie especially, given its announced end-of-life.

TIER 3 Differentiation amplifiers — where we win against incident.io and Rootly

Microsoft and Teams environment. Both principal rivals are Slack-native, Rootly explicitly so. Native Teams support is a structural advantage in a segment they serve poorly.

Infrastructure-as-Code maturity. Terraform-standard teams self-select as strong fits and value the platform's API depth — four hundred-plus endpoints, an MCP server, and SDKs.

A preference for structure over speed. Buyers who prize codified process, clear ownership, and audit and compliance maturity over a fast Slack bot are naturally inclined toward FireHydrant.

> **A note on positioning.** *FireHydrant should not be presented as a like-for-like alert-routing replacement for PagerDuty. The advantage lies in consolidation, lifecycle coverage, catalogue context, noise reduction, and predictable pricing — not in raw routing configurability.*

III. The scoring matrix

Each parameter is scored from one to three. **Tier 1 is a gate: any single score of one disqualifies the account, irrespective of the total.** Tiers 2 and 3 sum into a fit score and a winnability score respectively.

	1 POOR	2 ADEQUATE	3 STRONG
TIER 1 *Non-negotiable gate — any score of 1 disqualifies*			
Cost of downtime	No articulable cost; downtime an inconvenience	Hurts productivity or reputation; no hard number	Quantifiable revenue-per-minute or SLA penalties
Software-driven business	Non-digital, or pure IT-helpdesk buyer	Software-enabled, but not the core product	Always-on digital service at the core
Reliability function	No SRE / Platform / DevOps; ad-hoc response	Informal ownership; shared part-time pager	Dedicated team owning production incidents
Observability stack	No monitoring, or legacy and unknown	One basic tool; limited coverage	Datadog / Grafana / Prometheus in active use
TIER 2 *Fit score — sum, maximum 15*			
Maturity & trigger	No appetite to formalise	Aware of the need; no catalyst	Recent outage or new senior reliability hire
Org shape	Flat org; no specialised roles	Reliability ownership, but tribal	Dedicated roles; multi-team on-call; IC role
Company & team size	Under 500 or over 10,000 employees	500–10,000 employees; under 30 or over 200 engineers	500–10,000 employees; 30–200 engineers
Architecture	Single monolith; few services	Mixed; some decomposition	Complex microservice estate
Legacy alerting	No tool, and no plan to adopt one	Has alerting; no migration appetite	PagerDuty / Opsgenie deployed, with pain
TIER 3 *Winnability score — sum, maximum 9*			
Teams environment	Deeply embedded Slack shop	Mixed; no strong preference	Microsoft / Teams environment
IaC maturity	Manual configuration	Some automation; no Terraform standard	Terraform-standard; values API depth
Structure over speed	Wants the simplest Slack bot	Balanced	Prizes process, ownership, compliance

IV. Interpreting the result

FIT SCORE *Tier 2, maximum 15*

12 – 15	Strong fit. Prioritise.
8 – 11	Workable. Develop the weaker parameters before committing resource.
5 – 7	Marginal. Deprioritise unless a Tier 3 edge is exceptional.

WINNABILITY SCORE *Tier 3, maximum 9*

7 – 9	A structural edge. Lead with it, Teams above all.
4 – 6	Competitive. Win on merits.
3	Rivals are well-suited here. Expect a contest.

> **In sum.** *The ideal account clears all four Tier 1 gates, scores twelve or above on fit, and seven or above on winnability. Particular attention is owed to accounts that pass Tiers 1 and 2 yet score only three on winnability: attractive on paper, but structurally favourable to incident.io. These should be identified early and approached with care.*